

Ceph Command Cheatsheet

General access

'ceph -s', 'ceph health {detail}' Display cluster status;
 add watch cmd for real time
'ceph -w' Display running cluster status
'ceph report' Generate a comprehensive report of ceph
 conf
'ceph status -f json | jq' Output ceph status in json
 format
**'ceph --watch-debug|watch-info|watch-sec|watch-
 warn|watch-error'** watch debug/ info/ security /
 warning /error events
'ceph df' show cluster utilization statistics
'ceph fs ls' List file systems; other cmd "new|reset|rm"
'rados df {detail}' Get rados statistics
'ceph node ls' list basic info of nodes (json)
'ceph {-a} start|stop [mon|osd|mds|ceph-rados]' start
 or stop ceph services; '-a': {daemon_type}.id; similar
 usage: 'service ceph -a start|stop'
'ceph daemon {daemon_type}.id config show' Show
 config info (from the corresponding node); use "get|set
 config_parameter" see also "--admin-daemon"
**'ceph tell {daemon_type}.id or *} injectargs --
 {config_parameter} {value}'** similar to above cmd; but
 from any node
'ceph version' display version info

OSD access

'ceph osd create' Create OSDs
**'ceph auth del osd.x', 'ceph osd crush rm osd.x', 'ceph
 osd rm osd.x'** Remove OSDs (x=OSD ID)
'ceph osd stat' see if all OSDs are running
'ceph osd tree' View the CRUSH map
'ceph osd repair x' Repair OSD x
'ceph osd scrub x' send a scrub cmd to OSD x
'ceph osd crush reweight {name} {weight}' Adjust an
 OSD's crush weight
'ceph osd crush remove {name}' remove OSD or bucket
'[start | stop] ceph-osd id=x' Start or stop a osd with ID
 x from out/down (not always working)
'ceph osd [out|in|down|lost] {x}' Mark OSD
 in/out/down
'ceph osd map {poolname} {object-name}' Find the
 object location
'ceph osd find x' Search for an OSD and its location in a
 crush map
'ceph osd getmap -o file' Write a copy of the most
 recent OSD map to a file
'ceph osd getcrushmap -o file' Write a copy of the crush
 map from the most recent OSD map to file
'ceph osd dump' Dump the OSD map; with grep to get
 more information, e.g., 'replicated size' to display per
 pool placement group and replication levels
'ceph osd map <pool-name> <object-name>' Find out
 where a specific object is or would be stored in system
'ceph osd pause| unpause' Pause or unpause IO to OSD
'ceph osd [unset | set] noout' Set or unset the osd for
 noout state for maintenance purpose

PG access

'ceph pg dump { -o {filename} --format=json |plain}'
 Check placement group stats

'ceph pg {poolnum}.y query' Query a particular
 placement group (y=PG ID)
**'ceph pg dump_stuck [unclean | inactive | stale |
 undersized | degraded]'** Identify stuck placement groups
'ceph pg map {poolnum}.y' Find OSD for a particular pg
'ceph pg scrub y' scrub a placement group
'ceph pg repair' {PG-ID} Repair PG-ID

Pool access

'ceph df' Display pool usage stats
'rados -p pool_name ls' List pieces of pool using rados &
 verify that object stored the object
'ceph osd lspools' List pools
'ceph osd pool create', 'ceph osd pool delete' Create or
 delete a storage pool
**'ceph osd pool create {pool-name} {pg-num} [{pgp-
 num}] [replicated] [crush-ruleset-name] [expected-
 num-objects]**
**ceph osd pool create {pool-name} {pg-num} {pgp-num}
 erasure [erasure-code-profile] [crush-ruleset-name]
 [expected_num_objects]'**
'ceph osd pool create ecpool 12 12 erasure' Create an
 EC pool named ecpool with pg=pgp=12
**'ceph osd pool set-quota {pool-name} [max_objects
 {obj-count}] [max_bytes {bytes}]'** Set pool quotas for
 the maximum number of bytes and/or the maximum
 number of objects per pool
**'ceph osd pool rename {current-pool-name} {new-pool-
 name}'** Rename a pool
**'ceph osd pool mksnap | rmsnap {pool-name} {snap-
 name}'** make or remove a snapshot
'ceph osd pool set {pool-name} {key} {value}' Set pool
 values [size, min_size, crush_ruleset etc], e.g.,
'ceph osd pool set repool size 3' Set replicas num as 3
**'ceph osd pool set pool_name pg_num 512' 'ceph osd
 pool set pool_name pgp_num 512'** Set pool placement
 group sizes
'ceph osd pool get {pool-name} {key}' Get pool values
'ceph osd pool get poolname pg_num' Find out total
 number of placement groups being used by pool
'ceph osd dump | grep -i poolname' Find out replication
 level being used by pool
**'ceph osd pool delete ecpool ecpool --yes-i-really-
 really-mean-it'** delete ecpool
'rados mkpool <Name>; rados rm pool <Name>' Add or
 remove a pool to the configuration
'ceph osd blacklist ls' Display blacklisted clients

RBD access

'rbd -p pool_name list' Display rbd images in a pool:
**'rbd --pool rbd snap create
 pool_name/image_name@snap_name'** Create rbd
 snapshot
'rbd map {image}' Map image to OS; after mount, able
 to use dd; together with 'rbd resize'
'rbd snap ls pool_name/image_name' Display rbd
 snapshots; together with 'snap create'
'rbd showmapped' Display which images are mapped
 via kernel

Ceph Command Cheatsheet

'**rbd create {image-name} --size {megabytes}**' Create a block device
 '**rbd ls**' See the block device by listing them
 '**rbd info --image image-name**' Show image information
 '**rbd map block-device-name**' Show map info of block device

MDS access

'**ceph tell mds.{mds-id} injectargs --{switch} {value} [--{switch} {value}]**' Change configuration parameters on a running mds.
 '**ceph mds stat**' Enable debug messages
 '**ceph fs ls**' Check the CephFS filesystem list
 '**ceph mds add_data_pool <pool>**' Add data pool
 '**ceph mds cluster_down|cluster_up**' take down or up cluster

Mon access

'**ceph mon stat**' Show monitor stats
 '**ceph mon dump**' Dump monitor state
 '**ceph [-m monhost] {command}**' Issue command to monitor, e.g., mon_status
 '**ceph quorum_status -f json-pretty**' check mon quorum status for troubleshooting
 '**ceph compact**' Compact monitor's leveldb storage
 '**ceph injectargs <injected_args> [<injected_args>...]**' Inject conf args into monitor

CRUSH access

'**ceph osd crush dump**' View the crush map
 '**ceph osd crush rule list**' View the crush map rules
 '**ceph osd crush rule dump <crush_rule_name>**' View the detailed crush rule
 '**ceph osd getcrushmap -o crushmapdump**' Obtain crush map
 '**crushtool -d crushmapdump -o crushmapdump-decompiled**' Decompile crush map
 '**crushtool -c crushmapdump-decompiled -o crushmapdump-compiled ceph osd**' Compile crush map
 '**setcrushmap -i crushmapdump-compiled**' Set crush map

EC profile

'**ceph osd erasure-code-profile set {name} **
 **[[{directory=directory}] **
 **[[{plugin=plugin}] ** **[ruleset-root={root}] **
 **[ruleset-failure-domain={bucket-type}] **
 [[{key=value} ...] [--force]' Set EC profile
 '**...\ plugin=jerasure k={data-chunks} m={coding-chunks} **
technique={reed_sol_van|reed_sol_r6_op|cauchy_orig
|cauchy_good|liberation|blauum_roth|liber8tion} \ ...'
 '**...\ plugin=isa technique={reed_sol_van|cauchy} **
[k={data-chunks}] [m={coding-chunks}] \...
 '**...\ plugin=lrc k={data-chunks} m={coding-chunks} **
l={locality} [ruleset-locality={bucket-type}] \...
 '**...\ plugin=shc [k={data-chunks}] [m={coding-chunks}] **

[c={durability-estimator}] \... 4 types of plug-ins
 '**ceph osd erasure-code-profile [get | rm] myprofile**' Get or remove a EC profile; similar cmd '**ls**'

Benchmark access

'**ceph tell osd.x bench [NUMER_OF_OBJECTS] [BYTES_PER_WRITE]**' Bench write a OSD
 '**ceph osd perf**' Display OSD commit and apply latency
 '**rados bench -p {pool-name} {seconds} [write | seq | rand] --no-cleanup**' Beachmark an EC pool named ecpool with 10 seconds of write access without clearup [-t 100 (100 threads; default 16); -b size]; similar for '**fiio**' without mount
 '**rados -p pool-name cleanup --prefix benchmark**' Cleanup the pool
 '**rbd bench-write block-device-name --io-total|io-threads|io-total|io-pattern <seq|rand>**' Bench write a block device mounted
 '**rados -p rbd load-gen **
**--num-objects 50 --min-object-size 4M --max-object-size 4M **
**--max-ops 16 --min-op-len 4M --max-op-len 4M **
--percent 5 --target-throughput 2000 --run-length 60' Generate (simulated heavy) load on a Ceph cluster
 '**fiio --size=10G --ioengine=rbd --invalidate=0 --direct=1 --numjobs=10 --rw=write --name=fiojob --blocksize_range=4K-512k --iodepth=1 --pool=bench --rbdname=fio-test**' fio RBD benchmark

Auth access

'**ceph auth list | import**' List or import cluster keys
 '**ceph auth add <entity> [<caps> [<caps>...]]**' Add authentication info for a particular entity
 '**ceph auth del|get|get-key|export <entity>**' Delete, get or export key
 '**ceph auth get-or-create|get-or-create-key <entity> [<caps> [<caps>...]]**' get or add authentication/key info for a particular entity
 '**ceph config-key del|exists|get|list|put**' Access configuration key